

# Non-linear Least Squares Problem

## Final Project in Numerical Optimization

Bastian Haase

### Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Aim of this Work . . . . .	2
<b>2</b>	<b>Test Problems</b>	<b>3</b>
2.1	Freudenstein's Problem . . . . .	3
2.1.1	Characteristics . . . . .	3
2.2	A Classical Problem . . . . .	4
2.2.1	Characteristics . . . . .	4
2.3	Nielsen's Problem . . . . .	4
2.3.1	Characteristics . . . . .	4
2.4	Trig Problem . . . . .	5
2.4.1	Characteristics . . . . .	5
<b>3</b>	<b>The Algorithms</b>	<b>5</b>
3.1	Standard Algorithms . . . . .	5
3.2	Approximated Newton . . . . .	6
3.3	Brown-Dennis . . . . .	6
3.4	Dennis-Gay-Welsh . . . . .	7
3.5	Fletcher-Xu . . . . .	8
3.5.1	Optimal Parameter . . . . .	9
<b>4</b>	<b>Test Results</b>	<b>11</b>
<b>5</b>	<b>Appendix: Detailed Test Results</b>	<b>12</b>

# 1 Introduction

Non-linear least squares problems form an important subcategory of numerical optimization. One of their most important applications lies in data fitting, which is, for instance, an important tool in statistics. A typical nonlinear least squares problem has the form  $f(x) = \frac{1}{2} \sum_{i=1}^m r_i^2(x)$  for smooth functions  $r_j : \mathbb{R}^n \rightarrow \mathbb{R}$ , which are called *residuals*. Let  $r = (r_j)_{j=1,\dots,m}^T$  denote the *residual vector* and  $J$  its Jacobian. A short computation shows

$$\nabla f(x) = J(x)^T r(x) \quad \text{and} \quad \nabla^2 f(x) = J(x)^T J(x) + \sum_{i=1}^m r_i(x) \nabla^2 r_i(x),$$

which reveals an interesting feature of these problems: By computing the Jacobian  $J$ , one already gets the first term of the Hessian of  $f$ . Often, the residuals are small or nearly linear, in these cases the term  $J^T J$  is a good approximation of the Hessian that can be used in Newton's method. This method is called Gauss-Newton and performs very well in a wide range of problems. However, one frequently occurs problems with large residuals at the initial guess and/or at the solution where the described approximation might not yield satisfying results. In this work, I want to describe algorithms that also approximate the second term and analyze how well these algorithms perform in the large residual case.

## 1.1 Aim of this Work

The aim is to implement and compare 3 different strategies to approximate  $S(x) = \sum_{i=1}^m r_i(x) \nabla^2 r_i(x)$  in a Newton method. We will test them against standard implementations of the Newton method, Gauss-Newton and BFGS. The first strategy relies on approximating  $\nabla^2 f(x)$  by finite differences and to use this to approximate  $S(x)$ . Furthermore, we will implement two strategies proposed by Brown-Dennis ([BD71]) and Dennis-Gay-Welsh ([DGW81]) which approximate and update the term  $S(x)$ . It should be noted that Dennis-Gay-Welsh propose their update strategy in combination with a very sophisticated trust-region algorithm. This algorithm does not just adjust the trust region but also decides between their new update strategy or a standard Gauss-Newton model. As this work focuses on the approximation of  $S(x)$ , I will not implement the trust-region algorithm. Note that such an algorithm could also be paired with any other update strategy tested here. Finally, we will implement an algorithm proposed by Fletcher-Xu ([FX87]). It is a hybrid algorithm of Gauss-Newton and BFGS using line-search. This method chooses between Gauss-Newton and BFGS steps between each iteration depending on a parameter which we will tweak for every test problem.




Figure 1: Surface Plot of  $\sqrt{\text{Freudenstein}}$

## 2 Test Problems

In this section, we will briefly describe the test problems and their characteristics. All problems chosen are theoretical problems designed to test non-linear least squares problems (cmp. [BD71]), so they suit this project. For all test problems, our termination criterion will be a gradient norm less than  $10^{-8}$ .

### 2.1 Freudenstein's Problem

Freudenstein's problem is given by the equation

$$f(x) = \frac{1}{2}((x_1 - 13 + ((5 - x_2)x_2 - 2)x_2)^2 + (x_1 - 29 + ((x_2 + 1)x_2 - 14)x_2)^2)$$

which has dimensions  $m = n = 2$ . It has a global minimum at  $(5.0, 4.0)$  with a residual of zero and a local minimum at approximately  $(11.4128, -0.8968)$  where the residual is roughly 24.4921. Our initial guess is  $x_0 = (6.0, 6.0)$  with a residual of 12025. The colors indicate their position in Figure 1.

#### 2.1.1 Characteristics

This problem has small dimensions and the residuals are linear in one variable and cubic in the other. Our starting point has a large residual, while the minima have very low residuals. It will be interesting to see, whether one of the algorithms will converge to the local minimum with non-zero residual.

Figure 1 shows us that there is a "bump" between both minima and that our starting point is at a point with a very high slop in the  $x_2$  direction.

## 2.2 A Classical Problem

The following problem is a long-known, popular, non-linear least squares problem. It is given by the function

$$f(x) = \frac{1}{2} \sum_{k=1}^{10} (e^{-0.1kx_1} - e^{-0.1kx_2} - x_3 (e^{-0.1k} - e^k))^2$$

which has zero residual whenever  $x_1 = x_2$  and  $x_3 = 0$ . There is an additional zero residual at  $x = (1, 10, 1)$ . The dimensions of this problem are  $n = 3$  and  $m = 10$ . We will use the starting point  $x_0 = (0, 10, 20)$  where the residual is approximately 1031.154.

### 2.2.1 Characteristics

In this problem, the residual at the starting point is large, yet the residual at the minimum is zero. This suggests that Gauss-Newton should perform very well once we are close enough to a solution. We expect that the specialized algorithms should outperform the Gauss-Newton at the beginning, while it is unclear whether their performance close to a minimum will be competitive. We will investigate whether all algorithms will converge to the same solution. The residuals are linear in one component, yet exponential in the other two components.

## 2.3 Nielsen's Problem

Nielsen's problem is a data-fitting problem. It is given by the data in Table 1

j	0	1	2	3	4	5	6	7	8	9
y	2.0	0.0	2/3	0.0	2/5	0.0	2/7	0.0	2/9	0.0

Table 1: Data underlying Nielsen's problem

and the functional relationship

$$r_j(x) = x_1 x_3^j + x_2 x_4^j$$

for  $j = 0, \dots, 9$ . So the dimensions of the problem are  $n = 4$  and  $m = 10$ . The minimum is achieved at  $(0.97754, 0.97754, -0.65140, 0.65140)$  with a residual of 0.03734. Our initial guesses will be  $x_0 = (1.0, 1.0, -0.75, 0.75)$  (*Nielsen1*) where the residual is also only 0.13484. and  $x_0 = (1.0, 2.0, -2.0, 1.0)$  (*Nielsen2*) with a residual of 174009.4725.

### 2.3.1 Characteristics

The Hessian is not linear in any component. It will be interesting to see how the performance changes from the low residual initial guess to the large residual initial guess.

## 2.4 Trig Problem

Consider the problem given by

$$f(x) = \frac{1}{2} \sum_{k=1}^{20} \left( (x_1 - 0.2kx_2 - e^{0.2k})^2 + (x_3 + x_4 \sin(0.2k) - \cos(0.2k))^2 \right)$$

with dimensions  $n = 4$  and  $m = 2$ . The minimum is approximately at  $x = (-11.594, 13.204, -0.40344, 0.23678)$  with a residual of approximately 85822. We will choose the starting point  $x_0 = (25.0, 5.0, -5.0, -1.0)$  where the residual is approximately 7926693.

### 2.4.1 Characteristics

Note that the residuals are quadratic in every component. This problem has a large residual even at the solution, which will challenge our algorithms in a different way than the preceding problems do. We would expect that the specialized large residual problems outperform the other algorithms the most here.

## 3 The Algorithms

In this section, we will describe all algorithms that we have tested. Our descriptions will be more detailed when discussing the problems designed for non-linear least squares problems with large residual.

### 3.1 Standard Algorithms

We will test our specialized algorithms against a couple of standard algorithms. As we can actually compute the Hessian matrices for all four problems, we will use a modified implementation of Newton's method that modifies the Hessian by adding an appropriate multiple of the identity matrix if Cholesky factorization fails. Newton's method does not converge globally, but it provides quadratic convergence locally. The method can obviously not be used when the Hessian is not available and should be avoided when the Hessian is costly to construct.

We will furthermore use a standard implementation of BFGS that updates an approximation of the inverse of the Hessian. This algorithm is a standard algorithm that is very popular for problems where the Hessian is unavailable or costly. But, note that this algorithm is more generic and designed to handle many different types of optimization problems. It does not take our special setting into account. We have guaranteed global convergence but can only expect superlinear convergence, even when being close to the solution.

Lastly, we will also test Gauss-Newton. We described this method in the beginning as the standard algorithm for non-linear least squares problem with a small residual. It takes the special structure of our problem into consideration and basically approximates our Hessian for free (given that we compute the

gradient anyway), but this approximation will be inaccurate when the residual is large and the residual functions are not linear. Henceforth, we expect this method to perform poorly when the residuals are big. This should be very prominent in the trig problem, while it could be less influential for the other problems where the residual at the solution is small.

### 3.2 Approximated Newton

We have implemented Newton's method where we use an approximation of the Hessian instead of the real Hessian. But, to take our scenario into account, we will not approximate the whole Hessian via finite differences, but instead only the Hessians of the residuals  $\nabla^2 r_j(x)$ . We expect that this method performs very similar to Newton's method for all of our problems, as this behavior was observed in homework 3 in this class.

One should note that this method is very costly in terms of gradient evaluations and might therefore be impractical in problems of high dimensions with complex gradients. We will subsequently refer to this method as *appNew*.

### 3.3 Brown-Dennis

We describe Algorithm 2.1 in [BD71]. The idea is to approximate  $\nabla^2 r_j$  and update the approximation after each iteration. Then, one applies a regular Newton step. The initial approximation is chosen to be done by finite differences, i.e. the first step of the Brown-Dennis method is the same as the approximated Newton method.

Subsequently, let  $B_{i,k}$  denote the approximation of  $\nabla^2 r_i(x_k)$ . Then, the update formula is given by

$$B_{i,k+1} = B_{i,k} + (\nabla r_i(x_{k+1})^T - \nabla r_i(x_k)^T - B_{i,k} \Delta x_k) \frac{(x_{k+1} - x_k)^T}{\|\Delta x_k\|_2^2}.$$

The second summand can be interpreted as follows: Recall that  $B_{i,k}$  is an approximation of  $\nabla^2 r_i(x)$ . Thus, the term inside the brackets resembles the error of the first Taylor polynomial of  $\nabla r_i(x)$  centered around  $x_k$  and evaluated at  $x_{k+1}$ . Thus, the whole second term is approximating this error in the direction of the last step. From this point of view, one also notices the resemblance to BFGS and the SR1 update.

Brown-Dennis prove the following theoretical result concerning convergence of this method.

**Theorem 3.1.** *Let  $x^*$  be a critical point and let  $\Omega$  be a convex neighborhood. Let  $K_i \geq 0$  be constants such that for every  $x \in \Omega$ , we have  $\|\nabla^2 r_i(x) - \nabla^2 r_i(x^*)\| \leq K_i \|x - x^*\|$  for  $i = 1, \dots, m$ .*

*Then, whenever  $\nabla^2 f(x^*)$  is nonsingular, there exist constants  $\epsilon, \delta > 0$  such that the algorithm converges for any  $x_0 \in \Omega$  with  $\|x_0 - x^*\| < \epsilon$  and  $\|B_{i,0} - \nabla^2 r_i(x_0)\| \leq \delta$ .*

*Furthermore, the convergence is at least quadratic when  $f(x^*) = 0$ .*

*Proof.* Compare Theorem 1 and 2 in [BD71].  $\square$

**Remark 3.1.** Brown and Dennis prove more convergence results that do not need the assumption that a solution exists. However, these results only give superlinear convergence and require assumptions on the accuracy of their approximation  $B_{ik}$ .

### 3.4 Dennis-Gay-Welsh

In their paper [DGW81], Dennis, Gay and Welsh propose the algorithm *NL2SOL* that is implemented in many software packages for numerical optimization. Contrary to the last update formula, this approach updates an approximation to  $S(x)$  instead of approximations to each  $\nabla^2 r_i(x)$ , so it requires less storage. The approximation is based on the secant condition. This update is also sized before updating, as this has proven to yield a more robust algorithm.

In their paper, they propose this update formula in combination with a sophisticated algorithm that uses two different update strategies and a trust region approach. To be more precise, it uses a quadratic model on an elliptic trust region and starts with the Gauss-Newton method. Then, during every iteration, the algorithm decides whether to switch to the other method (Gauss-Newton or their new update strategy described below) and whether to adjust the trust region. The algorithm always tries to switch the method before shrinking the trust region.

It should be noted that this approach does not depend on the two methods that one has chosen. One could for instance use this strategy with Gauss-Newton and the Brown-Dennis update. Therefore, we have chosen to just focus on the update strategy in this work and our algorithm will therefore differ than the one proposed in [DGW81].

Let us now describe the update. Dennis, Gay and Welsh's goal was to find an update strategy that was simple but recovered all basic properties of  $S_k$ . Let  $B_k$  denote the approximation for  $S_k$ . They start with  $B_k = 0$ , i.e. with Gauss Newton. As  $S_k$  is always symmetric, they want  $B_k$  to be symmetric as well. They do not require definiteness, as  $S_k$  is often indefinite. From the observation

$$\begin{aligned} S_{k+1}\Delta x_k &= \sum_{i=1}^m r_i(x_{k+1}) \nabla^2 r_i(x_{k+1}) \Delta x_k \\ &\approx \sum_{i=1}^m r_i(x_{k+1}) (\nabla r_i(x_{k+1}) - \nabla r_i(x_k)) \\ &= J_{k+1}^T r_{k+1} - J_k^T r_{k+1} =: y_k \end{aligned}$$

we deduce the requirement  $B_{k+1}\Delta x_k = y_k$ . This is analogous to the derivation of the BFGS update. So, we have narrowed down our candidates for  $B_{k+1}$  to the set  $Q = \{S \mid S^T = S \text{ and } S\Delta x_k = y_k\}$ . The next theorem will help us make a good choice which element in this set should be used.

**Theorem 3.2.** *Let  $v \in \mathbb{R}^n$  be such that  $v^T \Delta x_k > 0$ . Then, for any positive definite symmetric matrix  $H$  for which  $H \Delta x_k = v$ , we get that*

$$B_{k+1} = \min_{S \in Q} \|S - S_k\|_{F,H} \quad (1)$$

where

$$B_{k+1} = B_k + \frac{(y_k - S_k \Delta x_k) v^T + v(y_k - S_k \Delta x_k)^T}{\Delta x_k^T v} + \frac{\Delta x_k^T (y_k - S_k \Delta x_k) v v^T}{(\Delta x_k^T v)^2}.$$

*Proof.* Compare theorem 3.1 in [DGW81].  $\square$

It remains to choose  $v$ . They suggest to choose  $v = \Delta \nabla f_k = J_{k+1}^T r_{k+1} - J_k^T r_k$ , thus choosing a norm induced by a matrix that sends  $\Delta x_k$  to  $\Delta(\nabla f_k)$ . Their rational is that this scaling should be similar to the scaling of the problem. Note that this choice is again analogous to the choice in BFGS.

One final ingredient is a scaling of  $B_k$  before updating it. This scaling is necessary to ensure that  $B_k$  converges to zero for zero residual problems. The update scalar is given by

$$\tau_k = \min \left\{ \frac{|\Delta x_k^T y_k|}{|\Delta x_k^T B_k \Delta x_k|}, 1 \right\}.$$

### 3.5 Fletcher-Xu

In their paper ([FX87]), Fletcher and Xu describe a hybrid algorithm that switches between regular Gauss-Newton steps and a modified BFGS step. We will use it in combination with a standard Armijo line-search with a starting step length of 1. This modified BFGS step has proven to be more efficient in least-squares problems. However, it might yield an indefinite update in which case the standard BFGS step is taken.

Let us first describe the modified BFGS method. If  $B_k$  is the current approximation to the Hessian, then, we obtain  $B_{k+1}$  via

$$B_{k+1} = B_k + \frac{\gamma \gamma^T}{\gamma^T \Delta x_k} - \frac{B_k \Delta x_k \Delta x_k^T B_k}{\Delta x_k^T B_k \Delta x_k}.$$

For the choice of  $\gamma$ , let us define

$$\begin{aligned} \gamma_1 &= \Delta(\nabla f_k) \\ \gamma_2 &= J_{k+1}^T J_{k+1} \Delta x_k + (J_{k+1} - J_k) r_{k+1} \end{aligned}$$

where  $\gamma_1$  yields standard BFGS and  $\gamma_2$  has proven to be more efficient in non-linear least squares problems. One should note that  $\gamma_1$  and  $\gamma_2$  are closely related,



which can be seen as follows (cmp. [ABF85]). Consider

$$\begin{aligned}
\gamma_1 &= \Delta(\nabla f_k) = \nabla^2 f_{k+1} \Delta x_k + o(\Delta x_k) \\
&= J_{k+1}^T J_{k+1} \Delta x_k + \sum_i r_i(x_{k+1}) \nabla^2 r_i(x_{k+1}) \Delta x_k + o(\Delta x_k) \\
&= J_{k+1}^T J_{k+1} \Delta x_k + (J_{k+1} - J_k) r(x_{k+1}) + o(\Delta x_k) \\
&= \gamma_2 + o(\Delta x_k)
\end{aligned}$$

where we used Taylor approximation twice. So,  $\gamma_2$  is just a variation of  $\gamma_1$  and usually performs better in non-linear least squares problems. Therefore, in our update, we choose  $\gamma_2$  unless

$$\Delta x_k^T \gamma_2 < 0.01 \Delta x_k^T \gamma_1$$

in which case we choose  $\gamma_1$  to maintain positive definiteness.

We have also tested this modified BFGS method by itself (without the hybrid component) under the name *MBFGS*.

The Fletcher-Xu algorithm is now very simple. Fix some  $\epsilon \in (0, 1)$ . Start with a Gauss-Newton step. Then, after every step, compute  $\frac{f_k - f_{k+1}}{f_k}$ . If the quantity is greater than  $\epsilon$ , continue with a Gauss-Newton step. Otherwise, apply the modified BFGS update.

If  $f^* = f(x^*) \neq 0$  and  $x_k \rightarrow x^*$ , then we get

$$\lim_{k \rightarrow \infty} \frac{f_k - f_{k+1}}{f_k} = 0$$

as the numerator converges to zero whereas the denominator does not. Hence, ultimately, BFGS steps will be taken. If  $f^* = 0$  and the order of convergence is superlinear, then Fletcher and Xu showed that Gauss-Newton steps will ultimately be taken. This behavior is of course desirable. They recommend to choose  $\epsilon = 0.2$  based on practical results, but we will approximate the optimal parameter for all problems ourselves in the next subsection.

Fletcher and Xu prove at least superlinear convergence in their paper under the assumption that the algorithm converges to a solution.

### 3.5.1 Optimal Parameter

In order to determine the optimal parameter  $\epsilon$ , we ran the algorithm for each problem with parameters  $0.05k$  for  $k = 1, \dots, 19$ . We will compare the number of iterations and the number of Gauss-Newton and BFGS steps.

The number of iterations was constantly 8 for Freudenstein's problem. However, for  $\gamma \leq 0.2$ , the algorithm took 5 Gauss-Newton and 3 BFGS steps, while it took 4 Gauss-Newton and 4 BFGS steps afterwards. We chose  $\epsilon = 0.2$  here although the difference is not relevant for a small problem like this one.

In the classical problem, the number of iterations was constantly 7, independent of  $\epsilon$ . Also, there were always 6 Gauss-Newton step and 1 BFGS step.

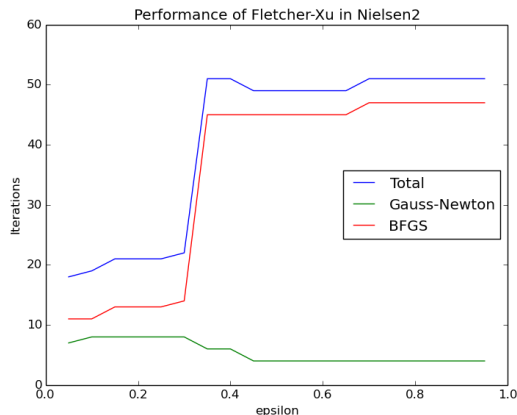


Figure 2: Iterations of Fletcher-Xu in Nielsen2

Nielsen1 was similarly unspectacular. The number of iterations is constantly 7, for big  $\epsilon$  this splits up into 3 Gauss-Newton and 4 BFGS steps. For small  $\epsilon \leq 0.3$ , the algorithm took 4 Gauss-Newton and 3 BFGS steps. We chose  $\epsilon = 0.2$

For Nielsen2 we can notice significant differences which are displayed in Figure 2. The optimal choice here is clearly  $\epsilon = 0.05$ . Lastly, the number of iterations for the Trig problem was constantly 27 with 1 Gauss-Newton and 26 BFGS steps. As this problem has a large residual nearly everywhere, this is not surprising.

Summarizing these results one sees that the choice was only significant in Nielsen2. We will later see that this problem is the most challenging one for all algorithms. The results for all other problems are to be expected and are on par with the theoretical properties of the algorithm. It can be concluded that the recommendation of  $\epsilon = 0.2$  is an effective choice.

## 4 Test Results

	Freudenstein	Classical	Nielsen1	Nielsen2	Trig
MNewton	9	16	6	21	9
Gauss-Newton	6	5	13	33	x
appNew	8	9	5	x	8
BFGS	17	29	8	40	29
MBFGS	10	11	7	54	26
Brown-Dennis	8	10	7	x	8
DGW	7	6	7	x	17
Fletcher-Xu	8	7	7	19	27

Table 2: Number of Iterations in all 5 Problems

The following Table 2 shows the number of iterations each method needed to converge. More details can be found in the appendix, which we will occasionally refer to in this section. In every table, x denotes a failure of the method. Note that BFGS converged to the strictly local minimum in Freudenstein’s problem, while all others converged to the global one. This behavior is certainly dependent on the line-search parameters, but it is still interesting to note that MBFGS with the same line-search converged to the global minimum.

Let us first compare Brown-Dennis, Fletcher-Xu and DGW. These are the algorithms specifically designed to solve large residual non-linear least squares problems. All three algorithms performed similarly well in Freudenstein, Classical and Nielsen1. While Fletcher-Xu was still able to converge in Nielsen2, the other two algorithms failed. This is surprising, especially considering that even Gauss-Newton converged here. But, in the Trig problem, all three algorithms converge and Brown-Dennis even outperforms Newton’s method. This is the only problem where the residual is not just large at the starting point, but also at the solution. The update strategy of DGW was slightly better than Brown-Dennis in all but the Trig problem. Maybe the trust-region-switch-model algorithm surrounding their update strategy would have improved the result here. But, in lights of the results in the Appendix, Brown-Dennis is much more costly than DGW or Fletcher-Xu. This is due to the fact that every  $\nabla^2 r_j$  gets updated individually.

It should be noted that Gauss-Newton performs very competitively for the first three if not four problems while being also very cheap to compute. It fails to converge in the last problem, so it seems to struggle with large residual at the solution, while it can handle large residuals at the starting point. MNewton and appNewton performed very well throughout, but while the first one requires Hessian constructions, the latter one approximates them with a huge number of gradient evaluations. This method also failed to converge in Nielsen2, which broke many algorithms.

Both BFGS and MBFGS obviously converged from every point. BFGS’s performance was mediocre, while MBFGS was competitive in the first three

problems. In fact, BFGS only beat MBFGS in Nielsen2. This shows that the modification do actually often yield better results in non-linear least squares problems. When looking at the appendix, one should consider the number of matrix factorizations with precautions: one could implement MBFGS just like BFGS in a way that the inverse of the Hessian gets approximated. In that case, no matrix factorizations are necessary. Note that the same cannot be said about Fletcher-Xu, as it also takes Gauss-Steps, which can only approximate the Hessian, but not its inverse. Therefore, one has to use a BFGS update formula for the Hessian here.

It is interesting to note that Fletcher-Xu seems to switch rather efficiently between Gauss-Newton and MBFGS despite its rather simple decision strategy.

Overall, the update strategy of Brown-Dennis often works well, especially in the case when the residue of the solution was large. But, on the other hand, it is very costly. DGW is cheaper, but performed poorly in the last problem which is very surprising. It seems as if a modified update strategy by itself is not the the answer to large residual non-linear least squares problems. Fletcher-Xu is very stable paired with decent performance and relatively cheap steps. It performs better than BFGS and MBFGS and is more stable than Gauss-Newton. Its steps are much cheaper than the steps in appNew and does not require the Hessian like MNewton. If it could be implemented without matrix factorizations, it would be the absolute clear winner, but even with this one flaw it is my general recommendation. One could also fix this flaw by transforming it into an iterative method. The BFGS update is of course very well suited for such an approach, the difficulty lies in handling the Gauss-Newton steps and the steps where the method switches. Nevertheless, such an approach would be feasible and could be implemented memory efficient.

## 5 Appendix: Detailed Test Results

In this appendix, you can find detailed test results. Here, FncEval stands for function evaluations and JEval stands for evaluations of  $J(x)$ . These results are analyzed in the section 4.

<b>Freudenstein</b>	FncEval	JEval	Matrix-Vector	Matrix-Matrix	Factorizations
MNewton	17	9	9	0	8
Gauss-Newton	13	13	0	8	6
appNew	0	29	8	0	8
BFGS	45	18	69	68	0
MBFGS	21	21	50	32	10
Brown-Dennis	0	21	32	33	8
DGW	0	8	57	57	7
Fletcher-Xu	15	15	15	15	7

Table 3: Detailed Results Freudenstein

<b>Classical</b>	FncEval	JEval	Matrix-Vector	Matrix-Matrix	Factorizations
Newton	37	16	16	0	23
Gauss-Newton	11	11	0	7	5
appNew	0	48	9	0	9
BFGS	63	30	117	116	0
MBFGS	23	23	55	35	11
Brown-Dennis	0	121	200	201	8
DGW	0	7	49	49	6
Fletcher-Xu	13	13	5	10	6

Table 4: Detailed Results Classical

<b>Nielsen1</b>	FncEval	JEval	Matrix-Vector	Matrix-Matrix	Factorizations
MNewton	11	6	6	0	5
Gauss-Newton	27	27	0	15	13
appNew	0	36	5	0	5
BFGS	18	9	33	32	0
MBFGS	15	15	35	23	7
Brown-Dennis	0	91	140	141	7
DGW	0	8	57	57	7
Fletcher-Xu	13	13	15	14	6

Table 5: Detailed Results Nielsen1

<b>Nielsen2</b>	FncEval	JEval	Matrix-Vector	Matrix-Matrix	Factorizations
MNewton	44	21	21	0	97
Gauss-Newton	69	67	0	35	33
appNew	x	x	x	x	x
BFGS	108	41	161	160	0
MBFGS	109	109	270	164	54
Brown-Dennis	x	x	x	x	x
DGW	x	x	x	x	x
Fletcher-Xu	47	37	55	45	18

Table 6: Detailed Results Nielsen2

<b>Trig</b>	FncEval	JEval	Matrix-Vector	Matrix-Matrix	Factorizations
MNewton	17	9	9	0	8
Gauss-Newton	x	x	x	x	x
appNew	0	65	8	0	8
BFGS	117	30	130	116	0
MBFGS	84	53	130	80	26
Brown-Dennis	0	201	320	321	8
DGW	0	18	137	137	17
Fletcher-Xu	85	53	130	8	27

Table 7: Detailed Results Trig

## References

- [ABF85] M Al-Baali and R Fletcher. Variational methods for Non-Linear Least Squares. *Journal of the Operational Research Society*, 36(May):405–421, 1985.
- [BD71] Kenneth M. Brown and John E. Dennis. New Computational Algorithms For Minimizing A Sum Of Squares Of Nonlinear Functions. *Research Report Number 71 - 6, Yale University*, 1971.
- [DGW81] John E. Dennis, David M. Gay, and Roy E. Welsh. An Adaptive Non-linear Least-Squares Algorithm. *ACM Transactions on Mathematical Software*, 7(3):348–368, 1981.
- [Fle86] R Fletcher. *Practical methods of optimization*. Wiley, 1986.
- [FX87] R Fletcher and C Xu. Hybrid Methods for Nonlinear Least Squares. pages 371–389, 1987.
- [GM78] Philip Gille and Walter Murray. Algorithms For The Solution Of The Nonlinear Least-Squares Problem. *SIAM J. NUMER. ANAL.*, 15(5):1425–1427, 1978.
- [NW04] Jorge Nocedal and Stephen J Wright. *Numerical Optimization*. Springer, 2004.